

CREDIT CARD FRAUD DETECTION

Import Dependencies

```
In [2]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score
```

The Dataset

```
In [3]: credit_card_data=pd.read_csv(r'C:\Users\HP\Downloads\archive (2)\creditcard.csv')
```

```
In [4]: credit_card_data.head()
```

```
Out[4]:
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	V22	V23
0	0.0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363787	...	-0.018307	0.277838	-0.110474
1	0.0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425	...	-0.225775	-0.638672	0.101288
2	1.0	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654	...	0.247998	0.771679	0.909412
3	1.0	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024	...	-0.108300	0.005274	-0.190321
4	2.0	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739	...	-0.009431	0.798278	-0.137458

5 rows × 31 columns

```
In [5]: credit_card_data.tail()
```

```
Out[5]:
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	V22	
284802	172786.0	-11.881118	10.071785	-9.834783	-2.066656	-5.364473	-2.606837	-4.918215	7.305334	1.914428	...	0.213454	0.111864	1
284803	172787.0	-0.732789	-0.055080	2.035030	-0.738589	0.868229	1.058415	0.024330	0.294869	0.584800	...	0.214205	0.924384	0
284804	172788.0	1.919565	-0.301254	-3.249640	-0.557828	2.630515	3.031260	-0.296827	0.708417	0.432454	...	0.232045	0.578229	-0
284805	172788.0	-0.240440	0.530483	0.702510	0.689799	-0.377961	0.623708	-0.686180	0.679145	0.392087	...	0.265245	0.800049	-0
284806	172792.0	-0.533413	-0.189733	0.703337	-0.506271	-0.012546	-0.649617	1.577006	-0.414650	0.486180	...	0.261057	0.643078	0

5 rows × 31 columns

Dataset Informations

```
In [6]: credit_card_data.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 284807 entries, 0 to 284806
Data columns (total 31 columns):
#   Column      Non-Null Count  Dtype
---  -
0    Time        284807 non-null  float64
1    V1          284807 non-null  float64
2    V2          284807 non-null  float64
3    V3          284807 non-null  float64
4    V4          284807 non-null  float64
5    V5          284807 non-null  float64
6    V6          284807 non-null  float64
7    V7          284807 non-null  float64
8    V8          284807 non-null  float64
9    V9          284807 non-null  float64
10   V10         284807 non-null  float64
11   V11         284807 non-null  float64
12   V12         284807 non-null  float64
13   V13         284807 non-null  float64
14   V14         284807 non-null  float64
15   V15         284807 non-null  float64
16   V16         284807 non-null  float64
17   V17         284807 non-null  float64
18   V18         284807 non-null  float64
19   V19         284807 non-null  float64
20   V20         284807 non-null  float64
21   V21         284807 non-null  float64
22   V22         284807 non-null  float64
23   V23         284807 non-null  float64
24   V24         284807 non-null  float64
25   V25         284807 non-null  float64
26   V26         284807 non-null  float64
27   V27         284807 non-null  float64
28   V28         284807 non-null  float64
29   Amount      284807 non-null  float64
30   Class       284807 non-null  int64
dtypes: float64(30), int64(1)
memory usage: 67.4 MB

```

```
In [7]: credit_card_data.isnull().sum()
```

```

Out[7]: Time        0
V1          0
V2          0
V3          0
V4          0
V5          0
V6          0
V7          0
V8          0
V9          0
V10         0
V11         0
V12         0
V13         0
V14         0
V15         0
V16         0
V17         0
V18         0
V19         0
V20         0
V21         0
V22         0
V23         0
V24         0
V25         0
V26         0
V27         0
V28         0
Amount      0
Class       0
dtype: int64

```

Hence it can be concluded that this given dataset has no null values.

```
In [8]: credit_card_data.shape
```

```

Out[8]: (284807, 31)

```

The given dataset has 284807 rows and 31 columns.

Distribution of legit transactions and fraudulent transactions

```
In [9]: credit_card_data.describe()
```

Out[9]:	Time	V1	V2	V3	V4	V5	V6	V7	V8
count	284807.000000	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05	2.848070e+05
mean	94813.859575	1.168375e-15	3.416908e-16	-1.379537e-15	2.074095e-15	9.604066e-16	1.487313e-15	-5.556467e-16	1.213481e-16
std	47488.145955	1.958696e+00	1.651309e+00	1.516255e+00	1.415869e+00	1.380247e+00	1.332271e+00	1.237094e+00	1.194353e+00
min	0.000000	-5.640751e+01	-7.271573e+01	-4.832559e+01	-5.683171e+00	-1.137433e+02	-2.616051e+01	-4.355724e+01	-7.321672e+01
25%	54201.500000	-9.203734e-01	-5.985499e-01	-8.903648e-01	-8.486401e-01	-6.915971e-01	-7.682956e-01	-5.540759e-01	-2.086297e-01
50%	84692.000000	1.810880e-02	6.548556e-02	1.798463e-01	-1.984653e-02	-5.433583e-02	-2.741871e-01	4.010308e-02	2.235804e-02
75%	139320.500000	1.315642e+00	8.037239e-01	1.027196e+00	7.433413e-01	6.119264e-01	3.985649e-01	5.704361e-01	3.273459e-01
max	172792.000000	2.454930e+00	2.205773e+01	9.382558e+00	1.687534e+01	3.480167e+01	7.330163e+01	1.205895e+02	2.000721e+01

8 rows × 31 columns

```
In [10]: credit_card_data['Class'].value_counts()
```

```
Out[10]: 0    284315
         1     492
         Name: Class, dtype: int64
```

0 ----> Legit transactions 1 ----> Fraudulent transactions

It can be observed that 492 fraud transactions were done using the particular credit card.

This dataset is highly unbalanced.

Seperating data of legit transactions and data of fraudulent transactions for data analysis

```
In [11]: legit=credit_card_data[credit_card_data.Class==0]
         fraud=credit_card_data[credit_card_data.Class==1]
```

```
In [12]: print(legit.shape)
         print(fraud.shape)
```

```
(284315, 31)
(492, 31)
```

Statistical measures of the data

```
In [14]: legit.Amount.describe()
```

```
Out[14]: count    284315.000000
         mean      88.291022
         std      250.105092
         min       0.000000
         25%       5.650000
         50%      22.000000
         75%      77.050000
         max     25691.160000
         Name: Amount, dtype: float64
```

```
In [15]: fraud.Amount.describe()
```

```
Out[15]: count      492.000000
         mean     122.211321
         std     256.683288
         min      0.000000
         25%      1.000000
         50%      9.250000
         75%     105.890000
         max     2125.870000
         Name: Amount, dtype: float64
```

Compare the values for both transactions

```
In [16]: credit_card_data.groupby('Class').mean()
```

Out[16]:	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V20	V21
Class													
0	94838.202258	0.008258	-0.006271	0.012171	-0.007860	0.005453	0.002419	0.009637	-0.000987	0.004467	...	-0.000644	-0.001235
1	80746.806911	-4.771948	3.623778	-7.033281	4.542029	-3.151225	-1.397737	-5.568731	0.570636	-2.581123	...	0.372319	0.713588

2 rows × 30 columns

Under Sampling

Build a sample dataset containing similar distribution of legit transactions and the fraudulent transactions

Number of fraudulent transactions ----> 492

```
In [17]: legit_sample=legit.sample(n=492)
```

Concatenating two data frames

```
In [18]: new_dataset =pd.concat([legit_sample,fraud],axis=0)
```

```
In [19]: new_dataset.head()
```

Out[19]:

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	V22
216409	140464.0	-0.700138	2.095846	-2.057056	1.648525	1.139541	-1.327866	1.802385	0.022209	-1.164496	...	0.014722	0.099030
154448	101573.0	2.001199	-0.071216	-0.747014	0.204079	-0.009969	-0.442587	-0.269615	-0.098480	1.578474	...	-0.558587	-1.478569
249119	154259.0	2.010719	-1.943531	-1.572516	-1.673460	-0.896737	0.035297	-0.928835	-0.156121	-1.445864	...	0.308526	1.126986
250603	154991.0	-0.514317	1.281833	-0.236068	-0.632398	0.126544	-1.116174	0.636817	0.238145	-0.014978	...	-0.258763	-0.630457
90427	63015.0	0.921564	-0.391010	-0.426926	0.256354	-0.332200	-1.330426	0.766585	-0.498255	-0.293433	...	0.087153	-0.156378

5 rows × 31 columns

```
In [20]: new_dataset.tail()
```

Out[20]:

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	V22	
279863	169142.0	-1.927883	1.125653	-4.518331	1.749293	-1.566487	-2.010494	-0.882850	0.697211	-2.064945	...	0.778584	-0.319189	0.000000
280143	169347.0	1.378559	1.289381	-5.004247	1.411850	0.442581	-1.326536	-1.413170	0.248525	-1.127396	...	0.370612	0.028234	-0.000000
280149	169351.0	-0.676143	1.126366	-2.213700	0.468308	-1.120541	-0.003346	-2.234739	1.210158	-0.652250	...	0.751826	0.834108	0.000000
281144	169966.0	-3.113832	0.585864	-5.399730	1.817092	-0.840618	-2.943548	-2.208002	1.058733	-1.632333	...	0.583276	-0.269209	-0.000000
281674	170348.0	1.991976	0.158476	-2.583441	0.408670	1.151147	-0.096695	0.223050	-0.068384	0.577829	...	-0.164350	-0.295135	-0.000000

5 rows × 31 columns

```
In [21]: new_dataset['Class'].value_counts()
```

```
Out[21]: 0    492
1    492
Name: Class, dtype: int64
```

```
In [22]: new_dataset.groupby('Class').mean()
```

Out[22]:	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V20	V21
Class													
0	96836.014228	-0.020909	-0.040966	0.045129	-0.047551	-0.060705	0.017105	0.035887	-0.020976	-0.068252	...	0.010844	-0.029636
1	80746.806911	-4.771948	3.623778	-7.033281	4.542029	-3.151225	-1.397737	-5.568731	0.570636	-2.581123	...	0.372319	0.713588

2 rows × 30 columns

Splitting the data into Features & Targets

```
In [23]: X=new_dataset.drop(columns='Class',axis=1)
Y=new_dataset['Class']
```

```
In [25]: print(X)
```

	Time	V1	V2	V3	V4	V5	V6	\
216409	140464.0	-0.700138	2.095846	-2.057056	1.648525	1.139541	-1.327866	
154448	101573.0	2.001199	-0.071216	-0.747014	0.204079	-0.009969	-0.442587	
249119	154259.0	2.010719	-1.943531	-1.572516	-1.673460	-0.896737	0.035297	
250603	154991.0	-0.514317	1.281833	-0.236068	-0.632398	0.126544	-1.116174	
90427	63015.0	0.921564	-0.391010	-0.426926	0.256354	-0.332200	-1.330426	
...	
279863	169142.0	-1.927883	1.125653	-4.518331	1.749293	-1.566487	-2.010494	
280143	169347.0	1.378559	1.289381	-5.004247	1.411850	0.442581	-1.326536	
280149	169351.0	-0.676143	1.126366	-2.213700	0.468308	-1.120541	-0.003346	
281144	169966.0	-3.113832	0.585864	-5.399730	1.817092	-0.840618	-2.943548	
281674	170348.0	1.991976	0.158476	-2.583441	0.408670	1.151147	-0.096695	
	V7	V8	V9	...	V20	V21	V22	\
216409	1.802385	0.022209	-1.164496	...	0.397503	0.014722	0.099030	
154448	-0.269615	-0.098480	1.578474	...	-0.198748	-0.558587	-1.478569	
249119	-0.928835	-0.156121	-1.445864	...	-0.036273	0.308526	1.126986	
250603	0.636817	0.238145	-0.014978	...	-0.017241	-0.258763	-0.630457	
90427	0.766585	-0.498255	-0.293433	...	0.397107	0.087153	-0.156378	
...	
279863	-0.882850	0.697211	-2.064945	...	1.252967	0.778584	-0.319189	
280143	-1.413170	0.248525	-1.127396	...	0.226138	0.370612	0.028234	
280149	-2.234739	1.210158	-0.652250	...	0.247968	0.751826	0.834108	
281144	-2.208002	1.058733	-1.632333	...	0.306271	0.583276	-0.269209	
281674	0.223050	-0.068384	0.577829	...	-0.017652	-0.164350	-0.295135	
	V23	V24	V25	V26	V27	V28	Amount	
216409	-0.010420	-0.344222	0.132028	-0.377454	0.325363	0.281977	137.98	
154448	0.588439	0.645595	-0.770757	-0.074031	-0.105877	-0.056268	11.99	
249119	-0.320766	-1.083708	0.240181	0.285725	-0.021176	-0.050808	181.00	
250603	0.108532	-0.120517	-0.394992	0.154329	0.231417	0.088870	3.87	
90427	-0.278427	0.500507	0.501105	1.045782	-0.147562	0.028808	215.00	
...	
279863	0.639419	-0.294885	0.537503	0.788395	0.292680	0.147968	390.00	
280143	-0.145640	-0.081049	0.521875	0.739467	0.389152	0.186637	0.76	
280149	0.190944	0.032070	-0.739695	0.471111	0.385107	0.194361	77.89	
281144	-0.456108	-0.183659	-0.328168	0.606116	0.884876	-0.253700	245.00	
281674	-0.072173	-0.450261	0.313267	-0.289617	0.002988	-0.015309	42.53	

[984 rows x 30 columns]

```
In [26]: print(Y)
```

```
216409    0
154448    0
249119    0
250603    0
90427     0
...
279863    1
280143    1
280149    1
281144    1
281674    1
Name: Class, Length: 984, dtype: int64
```

Splitting the data into Training Data & Testing Data

```
In [31]: X_train,X_test,Y_train, Y_test = train_test_split(X,Y,test_size=0.2,stratify=Y,random_state=2)
```

```
In [32]: print(X.shape,X_train.shape,X_test.shape)
```

```
(984, 30) (787, 30) (197, 30)
```

Model Training

Logistic Regression

```
In [33]: model=LogisticRegression()
```

Training the Logistic Regression Model with Training Data

```
In [34]: model.fit(X_train,Y_train)
```

```
Out[34]: ▾ LogisticRegression
LogisticRegression()
```

Model Evaluation

Accuracy Score

Accuracy on training data

```
In [35]: X_train_prediction = model.predict(X_train)
```

```
In [36]: training_data_accuracy = accuracy_score(X_train_prediction,Y_train)
```

```
In [37]: print('Accuracy on Training Data:', training_data_accuracy)
```

Accuracy on Training Data: 0.9542566709021602

Accuracy on test data

```
In [39]: X_test_prediction = model.predict(X_test)
test_data_accuracy = accuracy_score(X_test_prediction,Y_test)
print('Accuracy Score on Test Data:',test_data_accuracy)
```

Accuracy Score on Test Data: 0.9238578680203046

```
In [ ]:
```

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js